

# Kubernetes Observability Architecture with Grafana, Prometheus, and Loki

**Author:** Arman Singh Kshatri

March 31, 2025

## 1 High-Level Design Decisions and Tradeoffs

This observability architecture leverages **Kubernetes** to deploy and manage key monitoring components:

- **Grafana:** Visualization and dashboarding.
- **Prometheus:** Metrics collection and storage.
- **Loki:** Log aggregation, storage, and querying.

### 1.1 Design Choices

- **Specialized Applications:** Prometheus for metrics, Loki for logs, and Grafana for visualization, ensuring clear separation of concerns.
- **Kubernetes-Native Approach:** Services are deployed within the same namespace for easier discovery. Application and Dashboard use LoadBalancer, while Prometheus and Loki are internal (ClusterIP).
- **Single Namespace:** I've Kept everything inside a single namespace for simple service discovery.
- **Structured Logging:** JSON-based logs with request tracing for better querying.
- **Storage Strategy:** PersistentVolumes ensure log, metrics and dashboards are saved across restarts and rebuilds.
- **Probing:** Liveness & Readiness Probes ensure unhealthy pods restart automatically.

### 1.2 Tradeoffs

- **Complexity vs. Capability:** A multi-component setup introduces operational overhead but provides a **scalable and automated** observability stack.
- **Storage vs. Retention:** Prioritizing high-value logs optimizes storage but may limit historical analysis.
- **Kustomize instead of Helm:** Instead of Helm charts, I used Kustomize to manage raw Kubernetes manifests. This gives fine-grained control without the extra abstraction of Helm. However Requires manual config updates, but is more flexible.

## 2 Proof of Solution

This architecture **automates monitoring and debugging**, addressing slow and inconsistent troubleshooting.

- **Automated Monitoring:** Prometheus scrapes API metrics, Loki collects logs, and Grafana unifies the data into a single dashboard.
- **Performance Tracking:** Monitors response time, error rates, Go Routines, memory usage per container, and network usage in terms of response times etc.

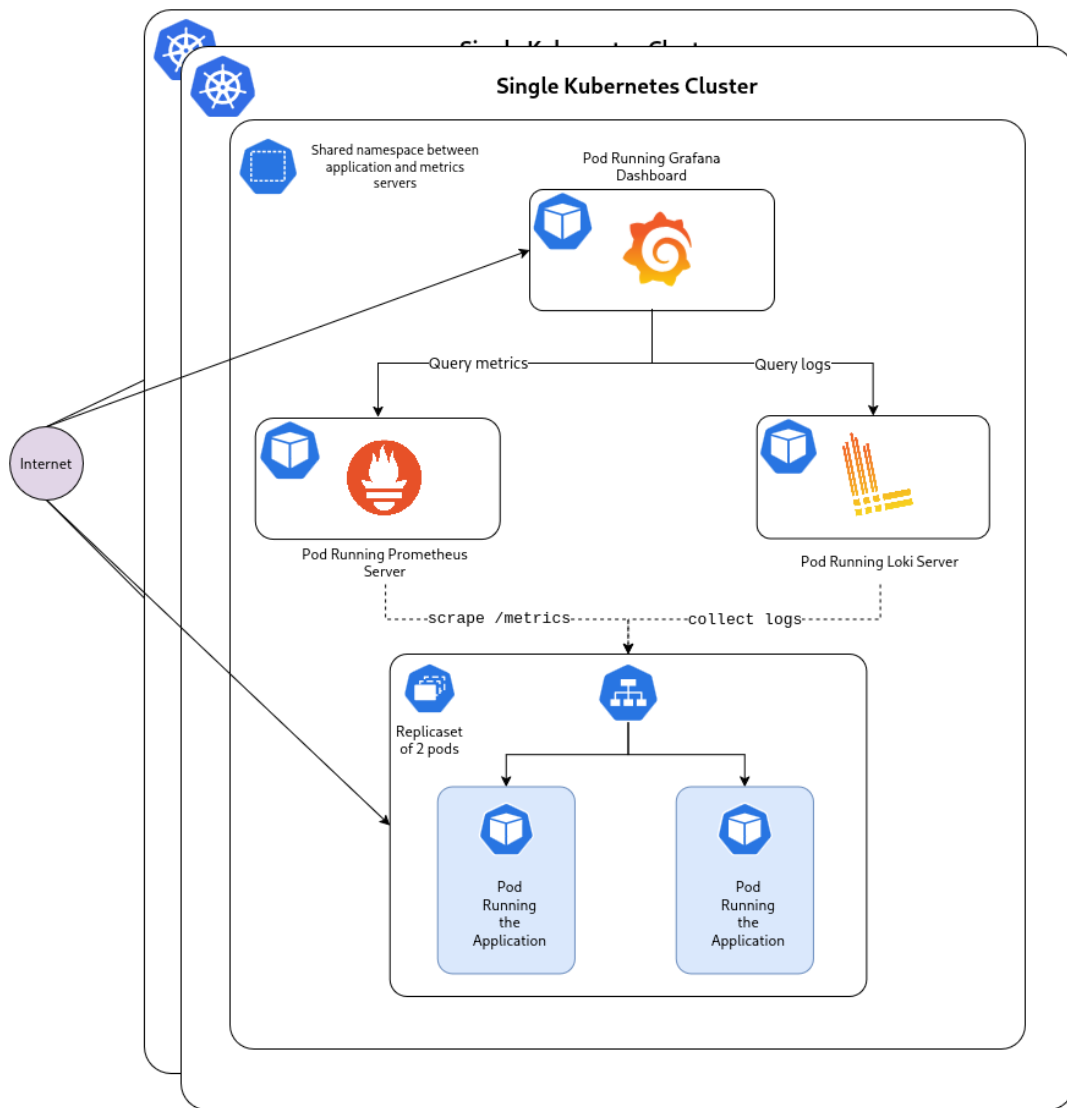


Figure 1: High-Level Design of Kubernetes Observability Architecture

### 3 Known Gaps and Justification

- **Limited Kubernetes Infrastructure Visibility:** Focuses on **application-level observability**; node-level monitoring is not a priority. Reason being Application has become a bottleneck, priority wise it should be handled first.
- **Historical Trend Analysis:** Limited retention(7d) since the focus is operational troubleshooting, However this is configurable and can be changed to more days.

### 4 Future Enhancements

- Setup Alerting Rules
- Prebuilt Dashboard Templates
- Improved Query Performance

### 5 Conclusion

This architecture **reduces MTTR**, enhances visibility, and **centralizes observability** for Kubernetes applications. By leveraging **Grafana, Prometheus, and Loki**, engineering teams gain an automated and scalable monitoring solution that minimizes manual debugging efforts.